

China Index Data Platform: End-to-End Pipeline on Snowflake, dbt, Airflow, and Streamlit

Author: Zhe Shen

1. Introduction

This document describes the design and implementation of the China Index Data Platform, a small end-to-end analytics warehouse for Chinese equity index daily prices.

The goal is to build a realistic but compact data pipeline that:

- Ingests daily OHLCV data for China A-share indices from AWS S3 bucket,
- Stores the data in Snowflake,
- Uses dbt to create clean staging, dimension, and fact tables,
- Uses Airflow to orchestrate a daily workflow, and
- Exposes the results in a simple Streamlit dashboard for analysis.

The main business questions are:

- How has a key index (or a small set of indices) performed over time in terms of cumulative returns and major drawdowns?
- How does volatility change over time, and when do we see calm vs. turbulent periods?
- What does the distribution of daily returns look like, and when do large positive or negative moves occur?
- Are there simple seasonality patterns by month or by weekday?

2. Data and Business Questions

2.1 Dataset

The original dataset contains daily OHLCV data (open, high, low, close, volume, turnover) for many China A-share indices over the period 2018-2023.

Each record corresponds to one index on one trading day and includes:

- trade_date - trading date,
- index_code - index identifier (e.g., SSE Composite, CSI 300),
- open, high, low, close - daily OHLC prices,
- volume - trading volume,
- turnover - trading value.

To keep the project scope and Snowflake storage manageable, the data is pre-filtered locally:

- Keep only a small set of major indices (e.g., SSE Composite Index and CSI 300),
- Restrict to a recent time range (e.g., 2019-2023),
- Keep only the columns needed for modeling.

The filtered result is saved as a single CSV file: index_daily_2019_2023_filtered.csv

This file is then uploaded to S3 and becomes the source for the end-to-end pipeline.

2.2 Business Questions

The platform is designed to support the following questions:

1. Performance and Drawdowns:

- How has each chosen index performed over time?
- What are the major drawdown periods and how deep were the losses?

2. Volatility and Extreme Days:

- How does rolling volatility evolve over time?
- When do we see large daily moves (big up days and big down days)?

3. Distribution of Returns:

- What is the shape of the daily return distribution?
- Are tails symmetric, and how common are large absolute moves?

4. Simple Seasonality:

- Are there clear differences in average returns by month?
- Are certain weekdays consistently stronger or weaker?

These questions motivate the design of the fact table and the metrics computed in dbt, as well as the visualizations in the Streamlit app.

3. Architecture Overview

The project follows a modern data stack pattern:

S3 → Fivetran → Snowflake → dbt → Airflow → Streamlit

3.1 High-Level Flow


1. S3 (Source Layer)

A pre-filtered CSV file containing index daily OHLCV data is stored in S3. Each row contains trade_date, index_code, open, high, low, close, volume, and turnover.

The screenshot shows the AWS S3 console interface. The left sidebar shows navigation options like 'Amazon S3', 'Buckets', 'Access management and security', 'Storage management and insights', and 'Account and organization settings'. The main area displays a bucket named 'vqy2214-mlds430-dw-project'. The 'Objects' tab is selected, showing a single item: 'index_daily_2019_2023_filter_ed.csv'. The object details show it's a CSV file, last modified on November 27, 2025, at 00:23:51 (UTC-06:00), with a size of 1.0 MB and a storage class of Standard. There are also buttons for 'Actions' (including Copy 53 URI, Copy URL, Download, Open, Delete, Create folder, and Upload).

2. Fivetran S3 → Snowflake (Raw Layer)

A Fivetran S3 connector ingests the CSV and loads it into Snowflake as a raw table in the schema MLDS430_KOALA_INDEX_RAW. Fivetran handles parsing, schema inference, and incremental loads.



You're now on the [Free Plan](#) ▾

Connections 5 Active At 3:32:54 AM

My Bookmarks Add connection

Try new connections for 14 days for free. [Learn more](#)

Search by name... Filter Save this view


Connection name	Source type	Destination	Status	Last synced
fivetran_metadata	Fivetran Platform	MLDS430Project	Active	2 hours ago
s3.mlds430project	Amazon S3	MLDS430Project	Active	2 hours ago

Our recommendations for you | Explore all 750 sources

- SQL Server Your tech stack [Add](#)
- Snowflake Your tech stack [Add](#)
- Google Drive Your tech stack [Add](#)
- Google Sheets Your tech stack [Add](#)
- SFTP Recommended [Add](#)
- Amazon RDS for PostgreSQL Recommended [Add](#)
- Workday Financial Management Recommended [Add](#)

3. Snowflake (Warehouse Layer)

The raw table in MLDS430_KOALA_INDEX_RAW is the single source of truth for dbt. No modeling is done here; Snowflake acts as the data warehouse and compute engine.



My Workspace

Search for files + Add new

Home 1.sql 2.sql HW3.sql test.sql Untitled.sql +

FIVETRAN_ROLE FIVETRAN_WAREHOUSE (X-Small) FIVETRAN_DATABASE MLDS430_KOALA_INDEX_RAW Share ...

My Workspace > Untitled.sql

```
1 --USE ROLE FIVETRAN_ROLE;
2 --USE DATABASE FIVETRAN_DATABASE;
3 --CREATE SCHEMA IF NOT EXISTS MLDS430_KOALA_INDEX_RAW;
4 --GRANT USAGE ON SCHEMA FIVETRAN_DATABASE.MLDS430_KOALA_INDEX_RAW TO ROLE FIVETRAN_ROLE;
5 --USE SCHEMA MLDS430_KOALA_INDEX_RAW;
6 -- CREATE OR REPLACE TABLE INDEX_DAILY_2019_2023_FILTERED AS SELECT * FROM
FIVETRAN_DATABASE.S3_INDEX_DAILY_2019_2023_FILTERED;
7 --select * FROM INDEX_DAILY_2019_2023_FILTERED WHERE TABLE_NAME = 1000001 ORDER BY DATE LIMIT 10;
8 SELECT * FROM STG_INDEX_DAILY limit 10;
```

Database Explorer

Objects Data Products

Search Filter

FIVETRAN_DATABASE DRAGON EARTHQUAKES_GEOCKO EARTHQUAKES_GEOCKO_PROJECT... FINAL_PROJECT_ANALYTICS.CHL... FINAL_PROJECT_CHIPMUNK FIVETRAN_ALPHABETICAL_HELP... FIVETRAN_ARISTOCRATIC_FRACT... FIVETRAN_BOBTAIL_JUSTIFIABLE... FIVETRAN_BYONE_EXCEEDED_S...

Results (5 hours ago)

TRADE_DATE	INDEX_CODE	OPEN	HIGH	LOW	CLOSE	VOLUME	AMOUNT	DATE_KEY	UNIQUE_KEY	
2019-06-07	1000039	40.0%	3299006	3240.95	3241.77	3197.91	3228.24	10185.12	19258559.6928	20210607 2399006_20210607
2019-01-28	1000039	30.0%	2520.03	2544.34	2499.35	2508.87	552.19	79259.016	20190128 1000039_20190128	
2019-04-19	1000001	30.0%	3250.15	3274.33	3226.81	3270.8	31519.8	31513192.0384	20190419 1000001_20190419	
2022-01-14	2399006	30.0%	3059.14	3128.25	3057.94	3119.41	15850.18	27146159.7184	20220114 2399006_20220114	
2019-09-09	1000039	30.0%	3845.44	3969.78	3816.96	3969.78	1712.49	3699919.2576	20190909 1000039_20190909	

4. dbt Transformations (Model Layer)

A dbt project (index_dbt) defines:

- stg_index_daily - a clean staging model,
- dim_date - a date dimension,
- fact_index_daily - a fact/mart table with returns and rolling metrics.

dbt also manages tests and basic documentation.

```
D:\MLDS430-Project-China-Index\index_dbt>dbt run
09:36:57  Running with dbt=1.10.15
09:36:58  Registered adapter: snowflake=1.10.3
09:37:00  Found 3 models, 8 data tests, 1 source, 494 macros
09:37:00  Concurrency: 4 threads (target='dev')
09:37:00
09:37:09  1 of 3 START sql table model MLDS430_KOALA_INDEX_RAW.dim_date ..... [RUN]
09:37:09  2 of 3 START sql table model MLDS430_KOALA_INDEX_RAW.stg_index_daily ..... [RUN]
09:37:13  1 of 3 OK created sql table model MLDS430_KOALA_INDEX_RAW.dim_date ..... [SUCCESS 1 in 3.41s]
09:37:13  2 of 3 OK created sql table model MLDS430_KOALA_INDEX_RAW.stg_index_daily ..... [SUCCESS 1 in 3.61s]
09:37:13  3 of 3 START sql table model MLDS430_KOALA_INDEX_RAW.fact_index_daily ..... [RUN]
09:37:16  3 of 3 OK created sql table model MLDS430_KOALA_INDEX_RAW.fact_index_daily ..... [SUCCESS 1 in 3.00s]
09:37:17
09:37:17  Finished running 3 table models in 0 hours 0 minutes and 17.42 seconds (17.42s).
09:37:17
09:37:17  Completed successfully
09:37:17
09:37:17  Done. PASS=3 WARN=0 ERROR=0 SKIP=0 NO-OP=0 TOTAL=3

D:\MLDS430-Project-China-Index\index_dbt>dbt test
09:37:42  Running with dbt=1.10.15
09:37:43  Registered adapter: snowflake=1.10.3
09:37:45  Found 3 models, 8 data tests, 1 source, 494 macros
09:37:45  Concurrency: 4 threads (target='dev')
09:37:45
09:37:47  1 of 8 START test not_null_dim_date_date_key ..... [RUN]
09:37:47  2 of 8 START test not_null_fact_index_daily_date_key ..... [RUN]
09:37:47  3 of 8 START test not_null_fact_index_daily_index_code ..... [RUN]
09:37:47  4 of 8 START test not_null_stg_index_daily_index_code ..... [RUN]
09:37:50  4 of 8 PASS not_null_stg_index_daily_index_code ..... [PASS in 2.61s]
09:37:50  1 of 8 PASS not_null_dim_date_date_key ..... [PASS in 2.62s]
09:37:50  5 of 8 START test not_null_stg_index_daily_trade_date ..... [RUN]
09:37:50  6 of 8 START test not_null_stg_index_daily_unique_key ..... [RUN]
09:37:50  2 of 8 PASS not_null_fact_index_daily_date_key ..... [PASS in 2.78s]
09:37:50  7 of 8 START test unique_dim_date_date_key ..... [RUN]
09:37:50  6 of 8 PASS not_null_stg_index_daily_unique_key ..... [PASS in 0.18s]
09:37:50  8 of 8 START test unique_stg_index_daily_unique_key ..... [RUN]
09:37:50  3 of 8 PASS not_null_fact_index_daily_index_code ..... [PASS in 2.84s]
09:37:50  5 of 8 PASS not_null_stg_index_daily_trade_date ..... [PASS in 0.29s]
09:37:50  8 of 8 PASS unique_stg_index_daily_unique_key ..... [PASS in 0.18s]
09:37:50  7 of 8 PASS unique_dim_date_date_key ..... [PASS in 0.29s]
09:37:51
09:37:51  Finished running 8 data tests in 0 hours 0 minutes and 6.39 seconds (6.39s).
09:37:51
09:37:51  Completed successfully
09:37:51
09:37:51  Done. PASS=8 WARN=0 ERROR=0 SKIP=0 NO-OP=0 TOTAL=8

D:\MLDS430-Project-China-Index\index_dbt>
```



Overview

Project **Database** **Group**

Tables and Views

- FIVETRAN_DATABASE
- MLDS430_KOALA_INDEX_RAW
- INDEX_DAILY_2019_2023_FILTERED
- dim_date
- fact_index_daily
- stg_index_daily

Overview

Project **Database** **Group**

Sources

- index_raw
- INDEX_DAILY_2019_2023_FILTERED

Projects

- index_dbt
- models
- dim
- dim_date
- fact
- fact_index_daily
- staging
- stg_index_daily

5. Airflow Orchestration

An Astro-based Airflow project defines a DAG index_etl_pipeline with three tasks: sync_fivetran, dbt_run, and dbt_test. The DAG can be scheduled daily to keep the models up to date.

The screenshot shows the Airflow UI for the index_etl_pipeline DAG. The DAG has three tasks: sync_fivetran, dbt_run, and dbt_test. All tasks are marked as successful. The dbt_run and dbt_test tasks were triggered manually, while sync_fivetran was triggered by PythonOperator.

Task ID	Map Index	State	Start Date	Try Number	Operator	Duration
dbt_test	✓ Success	Success	2025-11-30 03:43:09	1	BashOperator	00:00:07.759
dbt_run	✓ Success	Success	2025-11-30 03:42:49	1	BashOperator	00:00:18.369
sync_fivetran	✓ Success	Success	2025-11-30 03:42:47	1	PythonOperator	00:00:02.751

6. Streamlit Analytics

A Streamlit app connects to Snowflake, queries fact_index_daily, and provides interactive charts and tables for performance, drawdowns, volatility, extreme days, and seasonality.

Dashboard Filters

Index Selection

Select index

1000032 - SSE En...

Date Range

Start date

End date

Note: Data sourced from fact_index_daily table in Snowflake.

China Index Analytics Dashboard

Multi-index analytics for major Chinese equity indices: levels, cumulative return, drawdown, volatility and basic seasonality.

💡 Use the filters on the left to change indices and date range. All charts will update automatically.

Selected indices

SSE Composite Index SSE Energy

Performance & Drawdown

Volatility & Seasonality

Overall Index Performance

Index closing levels over time for the selected indices.



Cumulative Return

Cumulative return since the first date in the selected range.



Drawdown from Peak



Drawdown = Current cumulative return vs. historical peak (per index).

China Index Analytics Dashboard

Dashboard Filters

Index Selection

Select index

Date Range

Start date

2019/01/01

End date

2023/12/31

Note: Data sourced from fact_index_daily table in Snowflake.


Selected indices

SSE Composite Index SSE Energy

Performance & Drawdown Volatility & Seasonality

☁ 20-day Rolling Volatility

Standard deviation of daily returns over a 20-day rolling window.



🔝 Top 10 Up & Down Days


Top 10 Up Days

	trade_date	year	month	index_code	index_name	daily_return_%	close
1381	2020-07-06 00:00:00	2020	7	1000032	SSE Energy	6.29	1112.26
1532	2021-02-18 00:00:00	2021	2	1000032	SSE Energy	6.05	1163.79
1672	2021-09-09 00:00:00	2021	9	1000032	SSE Energy	5.74	1781.67
364	2020-07-06 00:00:00	2020	7	1000001	SSE Composite Index	5.71	3332.88
1674	2021-09-13 00:00:00	2021	9	1000032	SSE Energy	5.69	1822.08
33	2019-02-25 00:00:00	2019	2	1000001	SSE Composite Index	5.6	2961.28
1679	2021-09-22 00:00:00	2021	9	1000032	SSE Energy	5.39	1819.33
1683	2021-09-28 00:00:00	2021	9	1000032	SSE Energy	5.31	1761.85
1910	2022-09-05 00:00:00	2022	9	1000032	SSE Energy	5.29	2048.79
1692	2021-10-18 00:00:00	2021	10	1000032	SSE Energy	5.25	1710.44


Top 10 Down Days

	trade_date	year	month	index_code	index_name	daily_return_%	close
1277	2020-02-03 00:00:00	2020	2	1000032	SSE Energy	-8.44	1094.39
260	2020-02-03 00:00:00	2020	2	1000001	SSE Composite Index	-7.72	2746.61
1792	2022-03-15 00:00:00	2022	3	1000032	SSE Energy	-7.6	1469.8
1700	2021-10-28 00:00:00	2021	10	1000032	SSE Energy	-6.21	1466.82
1855	2022-06-20 00:00:00	2022	6	1000032	SSE Energy	-5.75	1767.27
79	2019-05-06 00:00:00	2019	5	1000001	SSE Composite Index	-5.58	2906.46
1684	2021-09-29 00:00:00	2021	9	1000032	SSE Energy	-5.57	1663.67
1694	2021-10-20 00:00:00	2021	10	1000032	SSE Energy	-5.55	1626.56
1689	2021-10-13 00:00:00	2021	10	1000032	SSE Energy	-5.38	1577.48
1867	2022-07-06 00:00:00	2022	7	1000032	SSE Energy	-5.38	1809.88


📊 Distribution of Daily Returns



■ Seasonality by Month (Average Daily Return)



▣ Seasonality by Weekday (Average Daily Return)



4. Ingestion Layer: S3 + Fivetran + Snowflake

4.1 S3 Source

The pre-filtered CSV file (e.g., index_daily_2019_2023_filtered.csv) is uploaded to an S3 bucket.

This file contains only the indices and date range needed for the project, which keeps the pipeline simple and efficient.

4.2 Fivetran S3 Connector

A Fivetran S3 connector is configured to:

- Read from the S3 bucket and path,
- Parse CSV files with a header row,
- Load them into Snowflake as a raw table.

The target configuration is:

- Database: the course Snowflake database,
- Schema: MLDS430_KOALA_INDEX_RAW,
- Table: a raw table such as INDEX_DAILY_2019_2023_FILTERED.

Fivetran can also be triggered via API in Airflow to force a manual sync, which is useful during development or demos.

4.3 Raw Table in Snowflake

After a Fivetran sync, Snowflake contains a raw table with one row per index_code and trade_date. The columns remain close to the original dataset. This table is referenced by dbt as a source and is not modified directly; all modeling and logic live in dbt models.

5. Transformation Layer: dbt Models

The dbt project index_dbt defines sources, models, tests, and configuration for Snowflake.

5.1 Source Definition

A dbt source is declared to reference the Fivetran raw table in MLDS430_KOALA_INDEX_RAW. This makes it easy to manage the connection and keeps SQL models focused on transformation rather than connection details.

5.2 stg_index_daily - Staging Model

Purpose: Clean and standardize the raw index data before further modeling.

Key tasks:

- Standardize column names to a consistent set: trade_date, index_code, open, high, low, close, volume, turnover.
- Cast trade_date to a proper date type.
- Cast all numeric columns to numeric types.

- Optionally filter to the subset of target indices (for example, SSE Composite and CSI 300).
- Create a date_key (YYYYMMDD as integer).
- Create a unique_key per index and date (e.g., concatenation of index_code and date_key).

Tests:

- not_null on trade_date, index_code, close, and unique_key.
- unique on unique_key.

This model is the clean base for all downstream metrics.

5.3 dim_date - Date Dimension

Purpose: Provide a calendar dimension for time-based analysis and seasonality.

Key tasks:

- Generate a continuous sequence of calendar dates covering the trading period (e.g., 2019–2023).
- Derive attributes such as:
 - year,
 - month,
 - day,
 - quarter,
 - weekday (e.g., Monday-Friday codes).
- Create a date_key (YYYYMMDD) that matches the one in stg_index_daily.

This table is used to filter by time period and to group returns by month or weekday for seasonality charts.

5.4 fact_index_daily - Fact / Mart Model

Purpose: Provide the main analytic table at the index-day grain.

Grain:

- One row per trading day per index.

Keys:

- date_key,
- index_code (and/or unique_key).

Measures and derived metrics:

- Base columns: open, high, low, close, volume, turnover.
- Daily return: Computed as $\text{close} / \text{lag(close)} - 1$ within each index_code, ordered by trade_date.

- Cumulative return: Constructed from cumulative log returns or a running product of $(1 + \text{daily_return})$ over time.
- Rolling 20-day average daily return: 20-day moving average of `daily_return` for each index.
- Rolling 20-day volatility: 20-day rolling standard deviation of `daily_return`.

All of these are implemented using SQL window functions (`lag`, `avg`, `stddev_samp` over partitions by `index_code` and ordered by `trade_date`). The model joins to `dim_date` on `date_key` to attach calendar attributes, enabling time-based grouping and filtering.

6. Orchestration Layer: Airflow DAG

The orchestration is implemented in an Astro-based Airflow project. The main DAG is named `index_etl_pipeline`.

6.1 DAG Tasks and Dependencies

The DAG contains three core tasks:

1. `sync_fivetran`
- A PythonOperator that reads environment variables(`FIVETRAN_API_KEY`, `FIVETRAN_API_SECRET`, `FIVETRAN_CONNECTOR_ID`) and calls the Fivetran REST API to force an S3 → Snowflake sync.
2. `dbt_run`
- A BashOperator that runs dbt run inside the Airflow container, targeting the `index_dbt` project.
3. `dbt_test`
- A BashOperator that runs dbt test after `dbt_run` completes.

The dependency chain is:

`sync_fivetran >> dbt_run >> dbt_test`

The DAG is scheduled to run daily, with `catchup=False` so that Airflow does not try to backfill historical dates. This reflects a realistic daily refresh in production while keeping the setup simple for the course project.

6.2 DAG Testing

A small test module (e.g., `tests/test_index_pipeline_dag.py`) checks:

- The DAG id and that all expected tasks exist,
- The dependency order between `sync_fivetran`, `dbt_run`, and `dbt_test`,
- The behavior of the Fivetran trigger function using mocks and environment variables.

These tests help ensure the DAG is correctly wired and make future changes safer.

7. Analytics Layer: Streamlit Dashboard

Instead of Tableau, the project uses Streamlit to build a lightweight, interactive analytics front end.

7.1 Connection and Filters

The app:

- Connects to Snowflake using the same private key configuration as dbt (SNOWFLAKE_PRIVATE_KEY_PATH, SNOWFLAKE_PRIVATE_KEY_PASSPHRASE and related settings).
- Queries the fact_index_daily table based on user selections.
- Provides a sidebar where the user can:
 - Select one or more index codes,
 - Choose a date range for analysis.

When the selections change, the app reloads the data from Snowflake into pandas and recomputes or re-displays the charts.

7.2 Visualizations

The Streamlit dashboard implements the visualization plan that was originally written for Tableau, adapted into interactive tabs:

1. Performance & Drawdown:
 - Line chart of index closing level over time,
 - Line chart of cumulative return over time,
 - Simple drawdown curve showing the percent drop from the running peak,
 - Optional summary metrics (e.g., cumulative return and max drawdown for the selected period).
2. Volatility & Extreme Days
 - Line chart of 20-day rolling volatility of daily returns,
 - Daily return chart, with visual emphasis on days where $|daily_return|$ exceeds a threshold (e.g., 3%),
 - Tables listing the top N up days and top N down days with dates and return values.
3. Seasonality
 - Charts showing average daily returns grouped by month and by weekday, using dim_date attributes,
 - Used to check for simple seasonal patterns in returns or volatility.

Together, these views allow the user to explore:

- Long-term performance and drawdowns,
- Changes in volatility over time,
- Extreme events (shock days),
- Basic seasonality patterns.

8. Implementation Notes and Environment

Key implementation details include:

- Authentication:
Snowflake authentication uses private keys, with paths and passphrases stored in environment variables for both dbt and Streamlit.
- Local development:
 - dbt is run locally via dbt run / dbt test,
 - Airflow is started via astro dev start in the Airflow/ directory,
 - Streamlit is started via streamlit run app.py in the streamlit/ directory.
- Performance:
The data volume (a few indices over a few years) is moderate, so Snowflake can easily handle transformations and queries. Window functions for returns and rolling volatility are executed in Snowflake rather than in Python, which keeps the pipeline efficient.

9. Conclusion

The China Index Data Platform demonstrates an end-to-end analytics pipeline for equity index data:

- Ingestion from S3 into Snowflake via Fivetran,
- Modeling with dbt staging, dimension, and fact tables,
- Orchestration with a daily Airflow DAG,
- Interactive analytics using a Streamlit dashboard.

The project shows how the tools from MLDS-430 course can be combined in a practical workflow to answer real business questions about performance, risk, and seasonality for Chinese stock indices. The design is simple but extensible, and it can serve as a template for larger, more complex data platforms in the future.